


Aiming for downtime

Cite as: Martin Paul Eve, "Aiming for downtime", *eve.gd: Martin Paul Eve*, January 02, 2022, <https://doi.org/10.59348/58q0h-49b49>.



Lots about contemporary computation stresses availability and *uptime*. It is important, for instance, that the OLH servers for which I have ultimate responsibility stay online. Otherwise, the articles that we publish would not be readable. As a result, various tools, such as “[Uptime Robot](#)”, have emerged that are designed to alert you when a server goes offline. There’s definitely something to be said about this demand for computational high-availability and the expectation that people be contactable/available in our digital world.

Indeed, I realized recently that this type of always-on thinking had saturated my personal computer usage. I have four NAS boxes, a headless server, and a desktop computer all running in my office. And I had become habituated, albeit subconsciously, to their instant availability and valorizing their uptime. These machines were always-on. It was convenient, quick, and easy.

This doesn’t seem like a very good situation given the dire hikes in energy prices that we face in the UK (not to mention global warming). So I am, instead, now embracing downtime. In fact, I’m even using Uptime Robot to monitor these machines... but with the explicit goal of minimizing their uptime. If I’m off for a 30 minute tea or lunch break, the machines are going off. They only take a few minutes to boot, after all.

The one exception is a Raspberry Pi unit that I use for our music system. This is staying on because it coordinates some of the NAS boxes to boot automatically, once per day, so that they can do their daily backup of all our systems. Overall, though, this is a step towards a greener (and cheaper)

computer system.